Collapse Around the Circuit: Induction Heads Survive Bounded Model Collapse While Their Errors Are Inherited

Matthew Purkiss (2697424)
University of the Witwatersrand, 1 Jan Smuts Ave,
Braamfontein, Johannesburg, 2000, South Africa

Sunay Master (2677874)
University of the Witwatersrand, 1 Jan Smuts Ave,
Braamfontein, Johannesburg, 2000, South Africa

Harshil Vallabh (2656158)
University of the Witwatersrand, 1 Jan Smuts Ave,
Braamfontein, Johannesburg, 2000, South Africa

Abstract

We study model collapse inside a mechanism we can read: the induction circuit of Singh et al. (2024). A two-layer attention-only transformer learns an in-context symbol task and is retrained for up to 20 generations on pools that mix a fraction λ of real data with data sampled from its predecessor. In the main sweep the circuit survives: test accuracy never falls below 0.979 and ablating layer 1 costs 0.66–0.76 accuracy in every generation. What degrades is the model's free generative choice. The probability it places on symbols absent from the context follows a two-parameter inheritance law, $m_{g+1} = a(1-\lambda)m_g + \varepsilon$ ($a = 0.951$, $R^2 = 0.994$), which predicts held-out real-data fractions. Removing those off-context samples from the synthetic pools cuts the gain a to 0.19, and training-data attribution finds them 12–23$\times$ over-represented among the most influential examples: the errors are carried by specific synthetic samples, not by a damaged circuit. A pre-registered extension of the law to higher sampling temperatures failed four of five tests. Collapse ran away only without real data, and the circuit then broke once correct labels in the training pool fell to chance level.

Keywords: model collapse; induction heads; in-context learning

Introduction

Language models adapt to patterns shown only in their prompt, and much of this in-context learning (ICL) is attributed to *induction heads*: a previous-token head copies each token's predecessor into its position, and an induction head one layer later attends from the current token to an earlier match and copies what followed it (Olsson et al., 2022). Singh et al. (2024) isolate this circuit in a minimal symbol-copying task. Separately, models trained recursively on generated data suffer *model collapse*, losing the tails of the data distribution (Shumailov et al., 2023, 2024); collapse is avoided when enough real data remains (Bertrand et al., 2023) or when data is accumulated rather than replaced (Gerstgrasser et al., 2024). These lines of work have not met. We ask whether recursive training damages the ICL circuit itself or only what the circuit produces, and we make a quantitative, pre-registered prediction about when collapse runs away.

Background Attention patterns suggest a circuit but are not causal evidence, so we follow Conmy et al. (2023) and ablate heads. To trace collapse back to specific training examples we use gradient-based data attribution (Pruthi et al., 2020), in the spirit of Chen et al. (2026).

Methods

Task and model. Each sequence gives $K{=}4$ symbols (from 16) a random label (from 6). A study block lists the pairs and a query block repeats the symbols in an independent order; the model predicts each label, which only content-based matching can solve. Unlike Singh et al. (2024) (two pairs, image symbols), we use tokens and query every symbol in shuffled order, so no positional shortcut works. The *base* model predicts only the query labels. The *extended* model then freely generates one more in-context symbol ($sq2$, uniform over the four in real data) and its label; $sq2$ is the degree of freedom that can collapse. The model has 2 attention-only layers of 4 heads ($d{=}64$, learned and rotary positions, tied embeddings, cosine learning-rate decay; 35k parameters). Hyper-parameters were chosen on validation loss and confirmed once on a held-out test seed.

Collapse pipeline. Generation 0 trains on a 20k-sequence real pool. Generation g trains a *fresh* model on a 20k pool with a fraction λ of real sequences and the rest sampled from generation $g{-}1$ at temperature $T{=}1$. Prompts stay real and all answers are model-generated (*prompt-real* regime). Each pool is split 90/10 into train and validation; test data is fresh real data. A generation must pass a strict gate (validation accuracy and a causal layer-1 check) or is retrained, which happened 20 times in 680 main-sweep generations with no final failures. We sweep $\lambda \in \{0, 0.1, 0.25, 0.5, 1\}$ for 12 generations, with 10 seeds for the extended model (20 generations at $\lambda{=}0$) and 5 for the base model.

Metrics. *Off-context mass*, the probability $p(sq2 \mid \text{context})$ places on symbols not in the context, measures hallucination; query NLL on real data measures confidence. The circuit is measured causally, as the accuracy lost when every layer-1 head is mean-ablated (at most about 0.83, down to chance at 1/6). This is agnostic to how the circuit is implemented, which matters because 5% of generations use a non-canonical variant that attention probes misread. We fit the inheritance law on $\lambda \in \{0, 0.25, 1\}$ and test it on $\{0.1, 0.5\}$. Interventions remove off-context or wrong-label samples from the synthetic pools, accumulate pools across generations, change the pool size, inject a one-off 5% label-noise pulse, and vary depth, width and task size.

Pre-registration. From one calibration seed, the law was fitted per temperature $T \in \{0.7, 1.3, 1.6\}$. It predicts a critical fraction $\lambda^* = 1 - 1/a(T)$ whenever $a(T) > 1$. Five pass/fail tests were committed to git before the confirmatory runs (5 seeds per cell).

Results

The circuit survives; its outputs do not. At generation 0 (test accuracy 0.998) a layer-0 head sends 0.79 of each label's attention to the preceding symbol and a layer-1 head 0.93 of each query's attention to its match. Accuracy then stays $\geq$0.979 in every generation, and the accuracy lost by ablating layer 1 is 0.71 at generation 0 and 0.69–0.74 at the last generation for every λ (Figure 1c). Confidence erodes instead: without real data, query NLL rises from 0.006 to 0.030 by generation 19 while accuracy is 0.996 (Figure 1b). The base model, with no free choice, shows only this effect: at $\lambda{=}0$ its accuracy stays 1.000 while query NLL rises from 0.0002 to 0.0058 by generation 11.

An inheritance law. Off-context mass grows from 0.013 to 0.138 by generation 19 without real data, and to 0.057, 0.035, 0.021 and 0.012 by generation 11 at $\lambda = 0.1, 0.25, 0.5$ and 1. The law fits all 630 transitions with $a = 0.951$ (95%

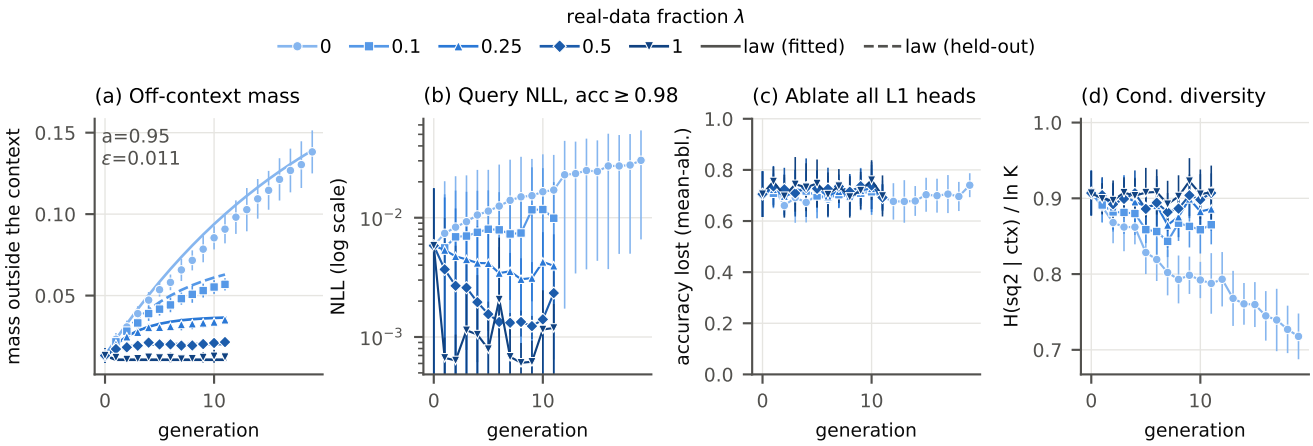


Figure 1: Main sweep (10 seeds, mean ± sd). (a) Off-context mass with the inheritance law; solid lines are fitted fractions, dashed lines are held-out predictions. (b) Query NLL rises without real data while accuracy stays ≥0.98. (c) Accuracy lost when all layer-1 heads are mean-ablated stays near 0.7: the circuit survives. (d) Conditional diversity of $sq2$ falls only without real data.

CI 0.942–0.959), $\varepsilon = 0.0105$ and $R^2 = 0.994$. Fitted without $\lambda \in \{0.1, 0.5\}$, it predicts them with RMSE 0.003, within one seed SD at 68% of generations (Figure 1a). It overshoots the plateau: for $\lambda{=}0$ it predicts $m^* = 0.21$, while mass is 0.13 at generations 16–19 and still rising. The gain holds across 3- and 4-layer, MLP and wider models ($a = 0.94$–1.04, $R^2 \geq 0.98$) and is lower on larger tasks (0.84 at $K{=}6$, $R^2 = 0.80$; 0.85 with 32 symbols); a one-layer model cannot learn the task (accuracy 0.64).

Specific synthetic samples carry the errors. Removing off-context samples from each synthetic pool cuts the gain from 0.95 to 0.19 (mass stays near ε; $R^2 = 0.05$), while removing wrong labels leaves it at 0.95. Attribution agrees: off-context samples are 12–23$\times$ over-represented among the top 1% most influential training examples. Accumulating pools lowers the gain to 0.62, and a 5k pool raises it above 1 (1.03 vs 0.91 for the matched 20k runs) and mass runs away to 0.58 without real data. A label-noise pulse decays at the predicted rate with real data (half-life 2.2 vs 2.1 generations at $\lambda{=}0.25$; 0.97 vs 0.93 at 0.5) but far more slowly without it (33 vs 14).

Table 1: Pre-registered temperature tests (confirmatory seeds 0–4). Only the qualitative prediction (test 5) held.

Test	Result	Observed ($T = 0.7, 1, 1.3, 1.6$)
1. $a(T)$ rises with T	fail	$a = -0.04, 0.93, 1.14, 1.04$
2. Inside 95% band	fail	3 of 12 cells
3. Contracting cells slow	fail	2 of 3 cells at $T{=}0.7$ (gain ≈ 0)
4. Runaway iff $a(1{-}\lambda) \geq 1$	fail	runaway only at $\lambda{=}0$; no seed in the 6 cells with $\lambda > 0$
5. Higher T: more mass, less homogenisation	pass	mass 0.011, 0.079, 0.77, 0.82; entropy 0.40, 0.78, 0.91, 0.93

The pre-registered prediction failed. Only the qualitative trade-off held (Table 1, test 5); the predicted critical fractions ($\lambda^* \approx 0.30$ at $T{=}1.3$, 0.52 at 1.6) do not exist: mass exceeded 0.5 in every seed at $\lambda{=}0$, yet no seed ran away in the six cells with 20–50% real data predicted to.

When collapse runs away, the circuit breaks. In those runs the share of correct labels in the synthetic pool fell to 0.15 ($T{=}1.3$) and 0.05 ($T{=}1.6$) by generation 9, but the circuit outlasted the labels. At $T{=}1.6$ it was intact at generation 4 with 26% correct labels (accuracy 0.93, layer-1 effect 0.67) and gone at generation 5 with 15% (0.37 and 0.06); at $T{=}1.3$ it broke at the same share (0.46 and 0.20). Both breaks came as correct labels fell to chance (1/6).

Discussion

In the prompt-real regime, bounded collapse leaves the induction circuit intact: every self-generated query is still answered by copying from a real context, which re-teaches the same computation. Collapse concentrates in the free choice, mainly as off-context hallucination inherited through the off-context samples. The circuit is lost only when copying stops paying, once correct answers are no more common than chance. The failed pre-registration bounds the linear account: at high temperature real data damps errors far more than a constant gain allows, so the law describes collapse locally but does not locate tipping points (its single-seed calibration also underestimated the gain at $T{=}1$: 0.86 vs 0.95). Survival also depends on supervising only the answers: with every token supervised (full language modelling) and no real data, the layer-1 effect falls from 0.83 to 0.71 and accuracy to 0.88 by generation 9, even with real prompts. Further limits are the 35k-parameter model, 16 symbols, 5–10 seeds and at most 20 generations. Two lessons carry over: accuracy can hide eroding confidence, and a modest share of real data, or accumulating data, sharply limits what a model inherits.

Contribution Statements (Not in Page Limit)

Group members: Sunay Master (2677874), Matthew Purkiss (2697424), Harshil Vallabh (2656158)

Project: *Interpretability of Model Collapse* (working title: *Collapse around the circuit: induction heads under recursive training*)

Research question. A small transformer learns the induction-head symbol task of Singh et al. (2024) and is then trained recursively on data generated by its previous generation (Shumailov et al., 2024). Does collapse degrade the induction circuit itself, or only what the circuit is asked to produce? How does the fraction of real data mixed into each generation change this, in both the base and the extended version of the task?

Division of work

Each member contributes one third (33.3%) of the project.

Member	Share	Experiments	Writing and admin
Sunay	33.3%	Core implementation and collapse pipeline; A+C sweeps and hyper-parameter selection; temperature test and poison-pulse runs	Method section; project outline and change plan
Matthew	33.3%	Single-generation analysis: training curves, attention maps, embeddings, convergence and loss floor; extra logging; headline figure	Introduction, Background and Results sections; extended abstract LaTeX; Faculty of Science AI ethics statement
Harshil	33.3%	Inheritance-law fit; gate audit and statistics; finite-data control run; mechanism analysis of the hallucination	Discussion and limitations; NeurIPS 2024 checklist; README.txt and requirements.txt

Sunay Master

- Build the study→query task generator (base and extended), the 2-layer transformer, the training loop, and the collapse pipeline (real/synthetic mixing, convergence gate with retries, collision-free seeding).
- Build the circuit probes (previous-token and induction scores, causal head ablation), conditional diversity and off-context mass; run the full sweep over real-data fractions, seeds and both modes; select hyper-parameters.
- Run the pre-registered temperature test of the critical real-data fraction and the poison-pulse experiment, and write the Method section.

Matthew Purkiss

- Produce the single-generation training curves, attention-map and embedding figures, and check that extended mode converges and reaches its predicted loss floor.
- Add validation loss and all-query perplexity to the logging, add mean-ablation to the probes, and build the composite headline figure.
- Write the extended abstract (Moodle template) including Introduction, Background and Results, and the Faculty of Science AI ethics statement.

Harshil Vallabh

- Fit the recursion model to the saved results, with held-out fractions as a test, and audit which generations fail a stricter convergence gate.
- Run the control that gives generation 0 the same data budget as later generations, and analyse which heads produce the off-context (hallucinated) probability.
- Write the Discussion and limitations, the NeurIPS 2024 checklist, README.txt and requirements.txt, and check that every number in the abstract traces to a file.

Shared by all members

- The two-page extended abstract (PDF), the full code, and checking every reference.
- The final contribution statement, which will report any changes to this plan and why.

Agreement

All group members have read and agree to this division of work.

Sunay Master Matthew Purkiss Harshil Vallabh

References

Bertrand, Q., Bose, A. J., Duplessis, A., Jiralerspong, M., & Gidel, G. (2023). On the stability of iterative retraining of generative models on their own data. *arXiv preprint arXiv:2310.00429*.

Chen, J., Luo, Y., & Pan, L. (2026). Mechanistic data attribution: Tracing the training origins of interpretable LLM units. *arXiv preprint arXiv:2601.21996*.

Conmy, A., Mavor-Parker, A., Lynch, A., Heimersheim, S., & Garriga-Alonso, A. (2023). Towards automated circuit discovery for mechanistic interpretability. *Advances in Neural Information Processing Systems*, *36*, 16318–16352.

Gerstgrasser, M., Schaeffer, R., Dey, A., Rafailov, R., Sleight, H., Hughes, J., ... Koyejo, S. (2024). Is model collapse inevitable? Breaking the curse of recursion by accumulating real and synthetic data. *arXiv preprint arXiv:2404.01413*.

Olsson, C., Elhage, N., Nanda, N., Joseph, N., DasSarma, N., Henighan, T., ... others (2022). In-context learning and induction heads. *arXiv preprint arXiv:2209.11895*.

Pruthi, G., Liu, F., Sundararajan, M., & Kale, S. (2020). Estimating training data influence by tracing gradient descent. In *Advances in Neural Information Processing Systems.*

Shumailov, I., Shumaylov, Z., Zhao, Y., Gal, Y., Papernot, N., & Anderson, R. (2023). The curse of recursion: Training on generated data makes models forget. *arXiv preprint arXiv:2305.17493*.

Shumailov, I., Shumaylov, Z., Zhao, Y., Papernot, N., Anderson, R., & Gal, Y. (2024). AI models collapse when trained on recursively generated data. *Nature*, *631*(8022), 755–759.

Singh, A. K., Moskovitz, T., Hill, F., Chan, S. C., & Saxe, A. M. (2024). What needs to go right for an induction head? A mechanistic study of in-context learning circuits and their formation. In *International Conference on Machine Learning* (pp. 45637–45662).

NeurIPS Paper Checklist

1. **Claims**

   Question: Do the main claims made in the abstract and introduction accurately reflect the paper's contributions and scope?

   Answer: [Yes]

   Justification: Every claim in the abstract is backed by a number in Results, and its scope is stated: the circuit survives only in the main sweep (bounded collapse) and breaks when collapse runs away, and the abstract reports that the pre-registered extension of the law failed four of five tests.

   Guidelines:

   - The answer NA means that the abstract and introduction do not include the claims made in the paper.
   - The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A No or NA answer to this question will not be perceived well by the reviewers.
   - The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
   - It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. **Limitations**

   Question: Does the paper discuss the limitations of the work performed by the authors?

   Answer: [Yes]

   Justification: The Discussion states that the circuit result holds only with real prompts and answer-only supervision (it degrades under full language-model training), and lists the 35k-parameter model, 16 symbols, 5–10 seeds and at most 20 generations. Results report where the inheritance law fails: it overshoots the $\lambda{=}0$ plateau and failed its pre-registered extension to other temperatures. The 2-page limit leaves no room for a separate Limitations section.

   Guidelines:

   - The answer NA means that the paper has no limitation while the answer No means that the paper has limitations, but those are not discussed in the paper.
   - The authors are encouraged to create a separate "Limitations" section in their paper.
   - The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
   - The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
   - The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
   - The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
   - If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
   - While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. **Theory Assumptions and Proofs**

   Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

   Answer: [NA]

   Justification: The paper has no theorems. The inheritance law $m_{g+1} = a(1-\lambda)m_g + \varepsilon$ is an empirical model fitted to data and tested on held-out real-data fractions, not a derived result.

Guidelines:

- The answer NA means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. **Experimental Result Reproducibility**

   Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

   Answer: [Yes]

   Justification: Methods give the task, architecture, pool size, real-data fractions, seeds, generations, sampling temperature, convergence gate and metrics. The remaining settings (AdamW, learning rate 10^{-3}, weight decay 0.01, batch 256, 200 warm-up steps then cosine decay, 6000 steps per extended and 4000 per base generation) are in `src/config.py` and `experiments/jobs_v2.py`. Runs use fixed seeds on CPU with one thread.

   Guidelines:

   - The answer NA means that the paper does not include experiments.
   - If the paper includes experiments, a No answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
   - If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
   - Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general. releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
   - While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
   (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
   (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
   (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
   (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. **Open access to data and code**

   Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

   Answer: [Yes]

Justification: The code, the raw outputs of all 311 runs and every analysis file are submitted with the paper and kept at `https://github.com/sunay777/NLP_Project1`. All data is synthetic and generated by `src/data.py`. `README.txt` gives one command per experiment and analysis, `requirements.txt` pins the packages (Python 3.12–3.13), and `python -m experiments.abstract_numbers` regenerates every number in the paper from the raw runs. Model weights are not included; they are rebuilt by rerunning the jobs.

Guidelines:

- The answer NA means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (`https://nips.cc/public/guides/CodeSubmissionPolicy`) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so "No" is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (`https://nips.cc/public/guides/CodeSubmissionPolicy`) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. **Experimental Setting/Details**

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer, etc.) necessary to understand the results?

Answer: [Yes]

Justification: Methods state the splits (each pool split 90/10 into train and validation; test data is fresh real data), that hyperparameters were chosen on validation loss and confirmed once on a held-out test seed, and the gate that retrains unconverged generations. The optimizer and exact values are in `src/config.py`, and the tuning results in `results/tuning_*.json`.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. **Experiment Statistical Significance**

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: Figure 1 shows mean $\pm$ one standard deviation over 10 seeds (stated in the caption); seeds change the initialisation, the real data and the sampling. The law's gain comes with a 95% bootstrap CI (1000 resamples of whole seed trajectories within each λ), and the held-out test reports RMSE and the share of generations within one seed SD. Two gaps: the intervention gains (scrubbing, accumulation, pool size, architecture and task size) are point estimates without CIs, and the symmetric SD bars in Figure 1b are clipped where they would go below zero on the log axis.

Guidelines:

- The answer NA means that the paper does not include experiments.

- The authors should answer ”Yes” if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g. negative error rates).
- If error bars are reported in tables or plots, The authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. **Experiments Compute Resources**

   Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

   Answer: [Yes]

   Justification: Given here, as the 2-page body has no room. All runs used CPUs only, one thread per run, on two Apple-silicon Macs; no GPU or cluster. With 9 runs in parallel, one extended generation took about 5 minutes per core. The 311 runs trained 3,293 generations (plus retries), about 330 CPU core-hours in total. Matthew’s share (200 runs) took about 17 hours of wall time on 9 parallel workers; Sunay’s share (111 runs, including the larger models) took about 26 hours on 9 workers over two sessions, about 5 hours of which were 4-layer runs repeated after a fix to the convergence gate. Earlier work not in the paper (a first sweep, hyper-parameter tuning, pilots and the temperature calibration) used additional compute.

   Guidelines:

   - The answer NA means that the paper does not include experiments.
   - The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
   - The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
   - The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn’t make it into the paper).

9. **Code Of Ethics**

   Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics `https://neurips.cc/public/EthicsGuidelines`?

   Answer: [Yes]

   Justification: The work uses only synthetic data, involves no human subjects or personal data, and releases a tiny model of no practical use outside the toy task. Our use of generative AI is declared in the Faculty of Science AI declaration attached to this report.

   Guidelines:

   - The answer NA means that the authors have not reviewed the NeurIPS Code of Ethics.
   - If the authors answer No, they should explain the special circumstances that require a deviation from the Code of Ethics.
   - The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. **Broader Impacts**

    Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [No]

Justification: The 2-page abstract has no room for this, so we give it here. This is foundational work on a toy task with no direct application. On the positive side, it supports keeping real data in training mixtures and monitoring confidence as well as accuracy, since models are increasingly trained on web data that contains model output. On the negative side, the label-noise pulse shows that injected errors persist far longer without real data, which is relevant to data poisoning; at this scale it gives an attacker nothing new.

Guidelines:

- The answer NA means that there is no societal impact of the work performed.
- If the authors answer NA or No, they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. **Safeguards**

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pretrained language models, image generators, or scraped datasets)?

Answer: [NA]

Justification: The only models are 35k-parameter transformers trained on a synthetic symbol task, and the only data is synthetic; neither carries a risk of misuse.

Guidelines:

- The answer NA means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. **Licenses for existing assets**

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: The task follows Singh et al. (2024), who are cited; we wrote our own implementation and used no existing datasets or models. The code uses PyTorch 2.14.1 (BSD-3-Clause), NumPy 2.5.4 (BSD-3-Clause) and Matplotlib 3.11.2 (Matplotlib license, based on the PSF license), within their terms.

Guidelines:

- The answer NA means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, `paperswithcode.com/datasets` has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. **New Assets**

    Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

    Answer: [Yes]

    Justification: The new assets are our code and the run outputs. `README.txt` and `RUN_V2.md` document how to run and analyse everything, `experiments/jobs_v2.py` lists every run, and each result file stores its full configuration. No human data is involved, so consent does not arise. The repository has no licence file.

    Guidelines:

    - The answer NA means that the paper does not release new assets.
    - Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
    - The paper should discuss whether and how consent was obtained from people whose asset is used.
    - At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. **Crowdsourcing and Research with Human Subjects**

    Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

    Answer: [NA]

    Justification: The work involves no crowdsourcing and no human subjects; all data is generated by a program.

    Guidelines:

    - The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.
    - Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
    - According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. **Institutional Review Board (IRB) Approvals or Equivalent for Research with Human Subjects**

    Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

    Answer: [NA]

    Justification: The work involves no human subjects or participants, so no ethics approval was needed.

Guidelines:

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.